

## *Plan*

### Last time

- power wall - difficult to shorten clock cycle
- memory wall - CPU speed versus memory latency
- analyzed  $A \cdot B$  on a single core with large slow and small fast memories.
  - **reuse** of local data needed for high performance
  - **blocking** helps to achieve reuse

### Today

- some terminology
- parallel computing models

## *Some loose terminology*

**PE:** any physical computational device

**Task:** is a set of instructions

**Stream:** is a set of instructions executed in sequence

**Paralel Task:** (parallel stream) is a set of tasks (streams) that are executed concurrently by multiple PEs

**Shared memory:** logical and physical addresses are the same for all PEs (global address space).

## *Some loose terminology*

**Communications:** data exchange among tasks.


**Distributed memory:** PEs can directly access only their own local data

- PEs must use explicit communication mechanism to access data they do not own and which is local to other PEs.

**Synchronization:** coordination of parallel tasks

- implemented by establishing a synchronization point where a task must wait until other tasks reach the same execution point.

## *Hardware models*



Several choices to make,

- hardware choices - # PEs, how powerfull, conectivity, etc.
- software choices

## *Parallel programming models*

### Software choices

- **Control**
  - How is parallel execution created?
- **Data sharing**
  - What is private and what is shared?
  - How is the data accessed?
- **Synchronization**
  - What operations can be used to coordinate parallelism?
  - What are the atomic (indivisible) operations?
- **Cost** of each of the above


## *Parallel machine models*

Flynn's taxonomy for instruction and data streams:

- **SISD** - Single Instruction, Single Data
- **SIMD** - Single Instruction, Multiple Data
- **MISD** - Multiple Instruction, Single Data (?)
- **MIMD** - Multiple Instruction, Multiple Data

## *SISD*

Classical Von Neumann machine:



One instruction stream is acted on by a single CPU


Single data is used as input at any clock cycle

Deterministic execution

Oldest type of a computer

(Later additions of pipelining and multiple functional units)

## *SIMD*



Each element of an array of data is acted upon by the same single instruction.




## *SIMD*

$P_1$		$P_2$		$P_3$	
load	$a(1)$	load	$a(2)$	load	$a(3)$
load	$b(1)$	load	$b(2)$	load	$b(3)$
$c(1)$	$= a(1) + b(1)$	$c(2)$	$= a(2) + b(2)$	$c(3)$	$= a(3) + b(3)$
store	$c(1)$	store	$c(2)$	store	$c(3)$

- PEs execute same instructions at same cycles (lockstep execution).
- PEs operate on different data sets organized into a common structure
- Specialized for problems with high degree of regularity.
- Connection Machine, MasPar, ILLIAC IV (by now long gone)
- GPUs (Nvidia, AMD), Cell (Sony Play Station), etc.

## *Shared memory MIMD*



Every PE may be executing a different instruction.

## Shared memory MIMD

$P_1$		$P_2$		$P_3$	
load	a(1)	load	a(2)	store	a(3)
load	b(1)	load	b(2)	load	a(2)
c(1)	= a(1) + b(1)	c(2)	= a(2) * b(2)	a(3)	= a(2) - b(3)
store	c(1)	c(2)	= c(1) + c(2)	store	a(3)

- Synchronous or asynchronous execution.
- SMPs, multi-core PCs.

## *Shared memory architecture*


- Common global address space visible to all PEs.
- Communication via shared data.
- Shared data needs to be protected by "locks" to avoid **race conditions**.
- Cache coherence mechanism required.

## *Shared memory threads model*

A program is divided into multiple concurrent streams (also called threads).

- **Main program** (master thread) asks the OS for all resources it needs to run.
- Master may perform some serial work, and then creates a number of tasks (slave threads) that the OS can run concurrently.
- Each thread may create new threads (its slaves).
- Each thread has local data but also shares resources of its masters.
- Threads communicate through global memory.
- Threads can be terminated and created until the master completes its job.

## *Shared memory threads model*



Things may go wrong if not treated properly.

### *Simple example - race condition*

global int s = 0; global int a[0:3] = {1,2,3,4}

thread 1	thread 2
for i = 0:1	for i = 2:3
s = s + a[i]	s = s + a[i]

- load and store are atomic, cannot be interrupted
- computations occur in local (to thread) registers

Sample trace:

Reg0 = a[0]	Reg0 = a[2]
Reg1 = s	Reg1 = s
Reg1 = Reg0 + Reg1	Reg1 = Reg0 + Reg1
<b>s = Reg1</b>	<b>s = Reg1</b>

## Shared memory

Reg0 = a[0]	Reg0 = a[2]
Reg1 = s	Reg1 = s
Reg1 = Reg0 + Reg1	Reg1 = Reg0 + Reg1
<b>s = Reg1</b>	<b>s = Reg1</b>
Reg0 = a[1]	Reg0 = a[3]
Reg1 = s	Reg1 = s
<div>Reg1 = a[0] <b>or</b> a[2]</div>	<div>Reg1 = a[2] <b>or</b> a[0]</div>
Reg1 = Reg0 + Reg1	Reg1 = Reg0 + Reg1
<b>s = Reg1</b>	<b>s = Reg1</b>

Need additional constructs to enforce correctness.



## *Shared memory*

### Advantages to programmer


- global address space
- communication by shared variables

### Disadvantages

- not scalable, more CPUs increase traffic on the memory-CPU path
- increased cache/memory management time
- programmer responsible for maintaining correct access of shared data

Expensive for a large number of processors.

## *Distributed memory MIMD*



- local memories only, no global address space
- local memory store has no effect on another memory (no cache coherency required)
- programmer explicitly controls how and when remote data is communicated (needs to know ids of all PEs).
- complex networks are used but can be as simple as Ethernet

## *Message passing MIMD*

Data is moved to and from PEs by sending and receiving **messages**,

- blocking - sender waits for acknowledgment
- or unblocking - can proceed with processing.

MPI is the standard for Message Passing Interface.

## *Distributed memory*

### Advantages:

- memory scale with number of PEs
- no overhead for maintaining cache coherency

### Disadvantages:

- programmer is responsible how and when data is communicated between PEs
- difficult to map efficiently existing data structures to distributed memory
- NUMA access times

## *More definitions*

**Granularity:** Task's computation to communication time ratio.

- **Coarse:** "large" amount of computational work is done between communication events - often distributed memory
- **Fine:** "small" amounts of computational work is done between communication events - often shared memory

Which one is better?

## *Performance metrics*

**Speed-up:** Ratio of sequential to parallel execution times with  $n$  PEs,

$$s_n = \frac{t_{seq}}{t_{par}}$$

Theoretically, ideal speed-up for  $n$  PEs is bounded by  $n$ .

Can  $s_n > n$  ?

## *Performance metrics*

$s_n > n$  in some situations like

- parallel task may follow different execution path
- cache effects
  - for large data set a single PE may miss on cache often
  - if there are many PEs with own caches misses occur less often

**Efficiency:** how well the resources are utilized,  $e_n = \frac{t_{seq}}{n \cdot t_{par}}$ .

## *Performance metrics*

**Scalability:** tells whether the algorithm can handle a growing amount of work or PEs efficiently.

We can increase either the number of processors, the size of the problem or both, and ask questions like

- can increase of the number of processors increase the speed-up?
- can increase of the problem size increase the speed-up?
- can a parallel system keep efficiency by increasing the number of processors and the problem size simultaneously?
- at what rate do we need to increase the problem size and the number of processors to maintain constant efficiency?



## *Amdahl's law*

Estimate of possible speed-up on  $n$  PEs.

### Assumptions

- a program consists of  $T$  "elementary" ops
- "elementary" ops all require one clock cycle to execute
- $\alpha \cdot T$  ops can be executed concurrently on  $n$  PEs
- $(1 - \alpha) \cdot T$  ops must be executed sequentially
- there is no overhead for memory access or communication

Derive upper bound on  $s_n$ .

## *Amdahl's law*

$$t_{seq} = T, \quad t_{par}(n) = \alpha \frac{T}{n} + (1 - \alpha)T$$
$$\lim_{n \rightarrow \infty} s_n = \lim_{n \rightarrow \infty} \frac{1}{\frac{\alpha}{n} + (1 - \alpha)} = \frac{1}{1 - \alpha}$$

Amdahl's Law: The speedup is bounded by the fraction of sequential code.

	$n = 2$	$n = 4$	$n = 16$
$\alpha = 0.5$	1.33	1.6	1.882
$\alpha = 0.7$	1.538	2.105	2.909
$\alpha = 0.9$	1.818	3.076	6.4

Need small  $\alpha$

- analogues to small miss rate for caches
- bad news ?

## *Chickens against ox*

When there is need to process larger jobs more resources must be provided.

One can either build a larger machine or add more small machines.

Which approach is better?

*If you were plowing a field, which would you rather use: Two strong oxen or 1024 chickens?*

Seymour Cray, the father of supercomputing.

## *Amdahl's law*

Consider program  $A$  with fraction  $\alpha$  of concurrent ops, and fraction  $(1 - \alpha)$  of sequential ops. Say  $A$  executes in time

- $t_{fast}$  on a fast PE  $P_{fast}$
- $t_{slow}$  on a slow PE  $P_{slow}$

and that  $\frac{t_{fast}}{t_{slow}} = f$ .

Say  $\alpha = 0.9$ .

For what value of  $f$  and  $n$ , is it better to use a system with slow PEs over a single fast PE?

## *Amdahl's law*

Possible speed-up with  $n$  slow PEs over the fast PE is

$$s_n = \frac{t_{fast}}{(1 - \alpha)t_{slow} + t_{slow}\frac{\alpha}{n}} = \frac{1}{\frac{(1-\alpha)}{f} + \frac{\alpha}{f \cdot n}}$$

$s_n$  is an increasing function of  $n$ , so

$$s_{n-1} < s_n < s_\infty = \frac{f}{1 - \alpha} = \frac{f}{0.1}$$

We have

$$s_n < 1, \quad \forall f = \frac{t_{fast}}{t_{slow}} < 0.1$$

so  $P_{fast}$  always better than  $n P_{slow}$  when  $t_{fast} < 0.1 \cdot t_{slow}$ .

*So why massively parallel machines?*

Parallel machines are for

- speeding-up a particular computation,
- solving larger and larger problems.

Invert the question:

How would a single PE perform if it were required to solve the problem solved by  $n$  PEs?

## *So why massively parallel machines?*

Assumptions:

- time  $t_{par}$  on multiprocessor with  $n$  PEs
- $\alpha$  fraction of time spent running on all PEs
- $(1 - \alpha)$  fraction of time running on a single PE

$$\begin{aligned} t_{seq} &= \alpha \cdot n \cdot t_{par} + (1 - \alpha) \cdot t_{par} \\ s_n &= \frac{t_{seq}}{t_{par}} = (1 - \alpha) + \alpha \cdot n \rightarrow \infty \end{aligned}$$

Good news ?

## *So why massively parallel machines?*

Assuming there is no communication cost we established that:

- Amdahls Law - speed-up is determined by the fraction  $\alpha$  of parallel code

$$s_n \leq \frac{1}{1 - \alpha}$$

- The choice between "oxens" and "chickens" depends on the ratio  $f = \frac{t_{fast}}{t_{slow}}$  and the value of  $\alpha$

$$s_n < \frac{f}{1 - \alpha}$$

- If  $n$  PEs execute in time  $t_{par}$  with  $\alpha$  fraction of parallel code then a single PE executes in time

$$t_{seq} = \alpha \cdot n \cdot t_{par} + (1 - \alpha) \cdot t_{par}$$



## *PRAM*

### Parallel Random Access Memory

- $P$  processors
- all memory access free
- all basic arithmetic operations cost 1 unit
- parallel step - number of ops performed simultaneously
- cost of a parallel algorithm is the number of parallel steps

Impossible to realize but

- good for analysis of algorithms
- provide upper bounds on possible speed-up

## *PRAM - binary summation*

$$s = \sum_{i=1}^{2^n} F_i$$

# flops ?


## *PRAM - binary summation*

$$s = \sum_{i=1}^{2^n} F_i$$

for  $j = 1 : n$

forall  $k = 1 : 2^{n-j}$

$$F[2^j k] = F[2^j k - 2^{j-1}] + F[2^j k];$$



# flops ?, speed-up  $s_N$  ?, speed-up  $s_P$  ?

## *Performance metrics*

For adding  $N = 2^n$  numbers on  $N/2$  PEs

$$E_N = \frac{t_{seq}}{N \cdot t_{par}} \approx \frac{1}{\log_2 N}$$

This process is not scalable.

For adding  $n$  numbers on  $p$  processors

$$E_p = \frac{n - 1}{p(\frac{n-p}{p} + \log_2 p)} \approx \frac{n}{n + p \log_2 p}$$

By choosing  $n = p \log_2 p$  we keep efficiency constant.

## *PRAM - parallel prefix*

$$s_k = \sum_{i=1}^k F_i, \quad k = 1, \dots, 2^n$$

$$s_1 = F_1$$

$$s_2 = F_1 + F_2$$

$$s_3 = F_1 + F_2 + F_3$$

$$\vdots$$

$$s_m = F_1 + F_2 + F_3 + \dots + F_m$$

# flops ?

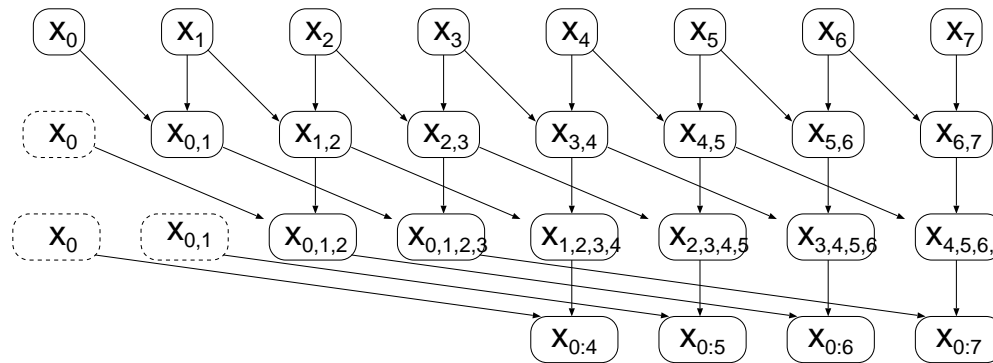
## *PRAM - parallel prefix*

$$s_k = \sum_{i=1}^k F_i, \quad k = 0, 1, \dots, 2^n$$

for  $j = 0 : n - 1$

forall  $k = 2^j : n$

$$F[j + 1, k] = F[j, k] + F[j, k - 2^j];$$



# flops ?, speed-up  $s_N$  ?, speed-up  $s_P$  ?

## *PRAM - parallel prefix*

Binary tree and prefix computation apply to any associative operation like:

- $F_0 \cdot F_1 \cdot F_2 \cdots F_{2^n-1}$
- $\max(F_0, F_1, F_2, \dots, F_{2^n-1})$

## *PRAM - linear recurrence*

Fibonacci sequence:

$$F_0 = 0, F_1 = 1,$$

$$F_n = F_{n-1} + F_{n-2}$$

Can it be parallelized?



## *PRAM - linear recurrence*

Another way of writing the recursion is

$$\begin{pmatrix} F_n \\ F_{n-1} \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} \begin{pmatrix} F_{n-1} \\ F_{n-2} \end{pmatrix}$$

Now use the parallel-prefix scheme.

Is it true for any

- linear recurrence?
- nonlinear?

*Next time*

Scheduling parallel tasks