

## *Plan*

- Domain decomposition applied to matrix operations
- Gaussian elimination

## *Domain decomposition for matrix computations*

### Domain Decomposition:

- the data associated with the problem is decomposed, tasks work on portions of the data
- common for matrix like operations
  - assign rows, columns or submatrices of a matrix to different PEs

### *PRAM matrix-matrix multiply*

$$C = A \cdot B \Rightarrow c(i, j) \leftarrow \sum_{k=1}^n a(i, k)b(k, j)$$

For a PRAM with "infinite" number of PEs:

- cost of a single "dot" product -  $\log_2 n$
- total cost -  $\log_2 n$ .

For a PRAM with  $P$  PEs:


- each PE computes  $\frac{n^2}{P}$  elements  $c(i, j)$
- total cost -  $2\frac{n^3}{P}$

## *Block matrix-matrix multiply*

Shared memory with  $P$  PEs, each with individual cache.

Split  $A, B, C$  into blocks of size  $b \times b$  (thus  $n = b \cdot p$ ).

$$C(i, j) \leftarrow C(i, j) + \sum_{k=1}^p A(i, k)B(k, j)$$



## *Block matrix-matrix multiply*

Assume  $P = p^2$   $C(i, j) \leftarrow C(i, j) + \sum_{k=1}^p A(i, k)B(k, j)$

for  $1 \leq i, j \leq p$

load block  $C(i, j)$  into fast memory

for  $k = 1 : p$

load block  $A(i, k)$  into fast memory

load block  $B(k, j)$  into fast memory

$C(i, j) \leftarrow C(i, j) + A(i, k)B(k, j)$

store  $C(i, j)$  into slow memory

Note that for all  $(i, j)$

- $A(i, k)$  is needed by all  $C(i, u)$ ,  $u = 1, \dots, p$
- $B(k, j)$  is needed by all  $C(v, j)$ ,  $v = 1, \dots, p$

### *All node shortest path*

- Bad: Simultaneous reads may create congestion on the interconnect.
- Good: only  $(i, j)$  modifies  $C(i, j)$  so no cache coherence is required.
- Good: no need for synchronization.
- Cost (time) will depend on the bandwidth of the interconnect between shared memory and caches.


## *All node shortest path*

Other algorithms have structure analogues to matrix-matrix multiply

- FloydWarshall algorithm for all pairs shortest paths.

### *All node shortest path*

$G = (V, E)$ ,  $V = \{1, 2, \dots, n\}$ .



Edge  $(i, j) \in E$  is assigned a non-negative distance  $w(i, j)$ .

Find shortest paths joining any pair of nodes.



## *All node shortest path*

Set  $W = \{w_{ij}\}$  where

$$w_{ij} = \begin{cases} 0 & i = j \\ w_{ij} & i \neq j, (i, j) \in E \\ \infty & i \neq j, (i, j) \notin E \end{cases}$$

Set  $D^{(m)} = \{d_{ij}^{(m)}\}$  where


- $d_{ij}^{(m)}$  shortest path from  $i$  to  $j$  containing at most  $m$  nodes

Assume  $D^{(m)}$  is known. We want to compute  $D^{(m+1)}$ .

When  $m = n - 1$  then  $D^{(m+1)}$  gives us all node shortest path.

How can we get  $D^{(m+1)}$  from  $D^{(m)}$  ?

### *All node shortest path*



Consider a new path from  $i$  to  $j$  through  $k$  with  $m + 1$  nodes.  
We have

$$d_{ij}^{(m+1)} = \min_{1 \leq k \leq n} (d_{ik}^{(m)} + w_{kj})$$

Note,  $d_{ik}^{(m)}$  can be  $\infty$ .

## *All node shortest path*

Notice the correspondence

matrix multiply  $\Leftrightarrow$  minimum sum

$$\text{L: } c_{ij} = \sum_{k=1}^n a_{ik} b_{kj} \quad \Leftrightarrow \quad \text{R: } d_{ij}^{(m+1)} = \min_{1 \leq k \leq n} (d_{ik}^{(m)} + w_{kj})$$

”.” on the left corresponds to ”+” on the right

”+” on the left corresponds to ”min” on the right

Computation of  $D^{(m+1)}$  can be viewed as multiplication of  $D^{(m+1)}$  with  $W$  with appropriate interpretation of elementary operations involved.

## *All node shortest path*


Cost ?

Note that


$$D^{(2)} = ?, \quad D^{(m)} = ?$$

Good matrix-matrix multiplication algorithm needed.

## *Matrix distributions*



## *Matrix distributions*



## *Triangular matrix multiply*


Matrices that are not "uniform".

$$y = \begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} t_{11} & & & & \\ t_{21} & t_{22} & & & \\ t_{31} & t_{32} & t_{33} & & \\ \vdots & & & \ddots & \\ t_{n1} & t_{n2} & \cdots & & t_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{pmatrix} = Tx$$

- How do we distribut work among PEs ?
- Do we need synchronization?

## *Triangular matrix multiply*

1D block-column distribution:



- load imbalance -  $P_4$  has the most work while  $P_1$  the least
  - about  $2n(n - p)$  ops for  $P_4$  and  $p^2$  for  $P_1$




## *Triangular matrix multiply*

- each  $P_i$  computes only partial "dot" products
  - $P_1$  computes  $\sum_{j=1}^{\frac{n}{p}} t_{ij}x_j, i = 1, \dots, n.$
- which  $P_i$  should accumulate partial "dot" products to get the final "dot" product?
- when should  $P_i$  start the final accumulation?

$P_i$  needs to be synchronized. Cache coherence needed?

## Triangular matrix multiply

1D cyclic distribution:




- $P_1$  performs about  $n + (n - p) + \dots + p = ?$  ops when computing  $x_{1+jp}^t, j = 1, \dots, \frac{n}{p}$

## *Triangular matrix multiply*

- load balance approximately maintained
- how are those modified columns added?
- we need to synchronize work
- cache coherence?

## *Triangular matrix multiply*


1D block-row distribution:



- $P_4$  computes  $\frac{n}{p}$  dot products of length  $\approx n$  for  $\frac{n}{p}(2n) = 2\frac{n^2}{p}$  ops.
- $P_1$  computes  $\frac{n}{p}$  dot products of length (about)  $\frac{n}{p}$  for  $\left(\frac{n}{p}\right)^2$  ops
- load imbalance
- no need for synchronization

## *Triangular matrix multiply*

1D cyclic distribution:



- no need to synchronize work
- load (almost) balanced

## *Solution of triangular systems*

$$T^{(0)}x = \begin{pmatrix} t_{11} & & & & \\ t_{21} & t_{22} & & & \\ t_{31} & t_{32} & t_{33} & & \\ \vdots & & & \ddots & \\ t_{n1} & t_{n2} & t_{n3} & \cdots & t_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ \vdots \\ b_n \end{pmatrix} = b^{(0)}$$

## *Solution of triangular systems*

$$x_1 = \frac{b_1}{t_{11}}, \quad b^{(1)} = b^{(0)} - t_{:,1} \cdot x_1, \quad \begin{pmatrix} t_{22} & & & \\ t_{32} & t_{33} & & \\ \vdots & & \ddots & \\ t_{n2} & t_{n3} & \cdots & t_{nn} \end{pmatrix} \begin{pmatrix} x_2 \\ x_3 \\ \vdots \\ x_n \end{pmatrix} = b_{2:n}^{(1)}$$

$$x_2 = \frac{b_2^{(1)}}{t_{22}}, \quad b^{(2)} = b^{(1)} - t_{:,2} \cdot x_2, \quad \begin{pmatrix} t_{33} & & & \\ t_{43} & t_{44} & & \\ \vdots & & \ddots & \\ t_{n3} & t_{n4} & \cdots & t_{nn} \end{pmatrix} \begin{pmatrix} x_3 \\ x_4 \\ \vdots \\ x_n \end{pmatrix} = b_{3:n}^{(2)}$$

How do we distribute work among PEs ?

## *Solution of triangular systems*


$$T^{(0)}x = \left( \begin{array}{cc|c} t_{11} & & \\ t_{21} & t_{22} & \\ \hline t_{31} & t_{32} & t_{33} \\ \vdots & & \ddots \\ t_{n1} & t_{n2} & t_{n3} & \cdots & t_{nn} \end{array} \right) \left( \begin{array}{c} x_1 \\ x_2 \\ \hline x_3 \\ \vdots \\ x_n \end{array} \right) = \left( \begin{array}{c} b_1 \\ b_2 \\ \hline b_3 \\ \vdots \\ b_n \end{array} \right) = b^{(0)}$$

- Get  $[x_1 \ x_2]^T$  (sequentially?) and
- modify  $b^{(0)}$  to  $b^{(2)}$  (here  $p = 2$ ) where all PEs can operate independently
- assign to each PE  $\frac{n-p}{p}$  rows of the "active" submatrices
- repeat until done



## *Solution of triangular systems*

Can we "pipeline" a bit?



- dynamic assignment of submatrices
- need to synchronize after each triangular solve and update

## *Solution of triangular systems*

What if the matrix is upper triangular ?

$$T^{(0)}x = \begin{pmatrix} t_{11} & t_{12} & t_{13} & \cdots & t_{1n} \\ & t_{22} & t_{23} & \cdots & t_{2n} \\ & & t_{33} & \cdots & t_{3n} \\ & & & \ddots & \vdots \\ & & & & t_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ \vdots \\ b_n \end{pmatrix} = b^{(0)}$$

## *Linear systems - Gaussian elimination*

- Model problem: Gaussian elimination

## *Linear systems - Gaussian elimination*

$$Ax = b$$

Gaussian elimination has two stages:

1. Transform the matrix of the system to upper triangular form
2. Solve the triangular system by backsubstitution

Issues to resolve:

- Distribution of work over PEs.
- Synchronization.

## *Solution of linear systems*

$$Ax = \begin{pmatrix} a_{11} & a_{12} & a_{13} & \cdots & a_{1n} \\ a_{21} & a_{22} & a_{23} & \cdots & a_{2n} \\ a_{31} & a_{32} & a_{33} & \cdots & a_{3n} \\ \vdots & \vdots & \vdots & & \vdots \\ a_{n1} & a_{n2} & a_{n3} & \cdots & a_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ \vdots \\ b_n \end{pmatrix}$$

Gaussian elimination: set  $i = 1$  and repeat:

1. divide row  $i$  of  $A^{(i)}$  and  $b^{(i)}$  by  $a_{ii}^{(i)}$
2. for  $j > i$  multiply new row  $\hat{a}_{i,:}^{(i)}$  by  $a_{ji}^{(i)}$  and subtract from row  $a_{j,:}^{(i)}$  to get  $a_{j,:}^{(i+1)}$
3. repeat (2) for the rhs to get  $b^{(i+1)}$
4. repeat (1) to (3) until  $i = n - 1$

## *Solution of linear systems*

$$\begin{pmatrix} \textcolor{red}{a}_{11} & a_{12} & a_{13} & a_{14} \\ \textcolor{green}{a}_{21} & a_{22} & a_{23} & a_{24} \\ \textcolor{green}{a}_{31} & a_{32} & a_{33} & a_{34} \\ \textcolor{green}{a}_{41} & a_{42} & a_{43} & a_{44} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \begin{pmatrix} \textcolor{blue}{b}_1 \\ b_2 \\ b_3 \\ b_4 \end{pmatrix}$$

1. Set  $i = 1$ , (here  $n = 4$ )
2. Divide row  $i$  of  $A$  and  $b$  by  $\textcolor{red}{a}_{ii}$  and store as  $\hat{a}_{i,i:n}, \hat{b}_i$ .
3. For rows  $j = i + 1, \dots, n$ , multiply  $(\hat{a}_{i,i:n}, \hat{b}_i)$  by  $\textcolor{green}{a}_{ji}$  and subtract from  $(\hat{a}_{j,i:n}, \hat{b}_j)$
4. Set  $i := i + 1$  and repeat from (2) until done.

## *Solution of linear systems*

After step 1:

$$\left( \begin{array}{c|ccc} a_{11} & a_{12} & a_{13} & a_{14} \\ \hline 0 & a_{22}^{(1)} & a_{23}^{(1)} & a_{24}^{(1)} \\ 0 & a_{32}^{(1)} & a_{33}^{(1)} & a_{34}^{(1)} \\ 0 & a_{42}^{(1)} & a_{43}^{(1)} & a_{44}^{(1)} \end{array} \right) \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \begin{pmatrix} b_1 \\ \hline b_2^{(1)} \\ b_3^{(1)} \\ b_4^{(1)} \end{pmatrix}$$

After step 2:

$$\left( \begin{array}{c|ccc} a_{11} & a_{12} & a_{13} & a_{14} \\ \hline 0 & a_{22}^{(1)} & a_{23}^{(1)} & a_{24}^{(1)} \\ 0 & & a_{33}^{(2)} & a_{34}^{(2)} \\ 0 & & a_{43}^{(2)} & a_{44}^{(2)} \end{array} \right) \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \begin{pmatrix} b_1 \\ \hline b_2^{(1)} \\ b_3^{(2)} \\ b_4^{(2)} \end{pmatrix}$$

## *Solution of linear systems*

After step  $n - 1$

$$\underbrace{\begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ 0 & a_{22}^{(1)} & a_{23}^{(1)} & a_{24}^{(1)} \\ 0 & 0 & a_{33}^{(2)} & a_{34}^{(2)} \\ 0 & 0 & 0 & a_{44}^{(3)} \end{pmatrix}}_U \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \underbrace{\begin{pmatrix} b_1 \\ b_2^{(1)} \\ b_3^{(2)} \\ b_4^{(3)} \end{pmatrix}}_d$$




## *Solution of linear systems*

### Questions:

- What if  $a_{ii}^{(i)} = 0$  ?
- What if  $|a_{ii}^{(i)}|$  is very small ?
- Pivoting: find  $k = \operatorname{argmax}_{j \geq i} |a_{ji}|$ .
- swap rows  $i$  and  $k$
- How do we distribute  $A$  among PEs ?

## *Block row distribution*


Distribute data.



- load imbalance
- need to synchronize (where ?)

## *Block row distribution*


Distribute data.



- cyclic distributes work more evenly
- need to synchronize (where ?)

## *Block row dynamic distribution*

Another option:



- Divide evenly amount of work left.
- At some point switch from multiple PEs to a single PE.

## *Solution of linear systems*

### Steps


1. local pivot (binary tree)
2. synchronize
3. global pivot
4. synchronize
5. get a copy of the pivot row (all or part ?)
6. update
7. synchronize
8. repeat

## *2D block cyclic distribution*

Many other ways to organize Gaussian elimination.

Each new architecture may need a new kind of reorganization.

A good project.



*Next time*

## Introduction to Pthreads