

## *Plan*

### Last time

- Domain decomposition for matrix-matrix multiplication
- All nodes shortest path by Floyd-Warshall
- Domain decomposition for triangular matrix-vector multiplication
- Domain decomposition for triangular solve
- Model problem: Gaussian elimination

### Today

- Intro to Pthreads
  - examples are taken from  
<https://computing.llnl.gov/tutorials/pthreads/>

## *Threads model*

A **program** is a collection of execution paths.

A **process** is composed of a sequence of instructions (its code), as well as input and output sets (its data).

- OS maintains information about process such as the program counter, registers, code segment, data segment, stack, heap, etc.

A **threads** is a sequence of instructions within a process that can be scheduled to run by the OS.

- Threads are within the same process address space, which they read and write **asynchronously**.

## *Threads model*

Threads maintain some private information such as

- pointers to individual stack to store private to the thread values
- private registers
- heap space, shared between threads belonging to the same process

The process itself is the main thread

- it first performs some serial work, and
- then creates a number of threads that are run by the OS.

## *Shared-memory programming model*

Threads **communicate** implicitly by writing and reading shared variables.

**Locks** assure that only one thread at a time updates the same global address within the shared memory.

## *POSIX threads*

NOTE: In what follows pthread codes (slightly modifeied) are adopted from <https://computing.llnl.gov/tutorials/pthreads/>

### Basic Pthreads types

---

<code>pthread_t</code>	a descriptor and ID
<code>pthread_mutex_t</code>	a lock for pthreads
<code>pthread_cond_t</code>	a conditional variable
<code>pthread_attr_t</code>	a descriptor for pthreads properties
<code>pthread_mutexattr_t</code>	a descriptor for mutex properties

## *POSIX threads*

### Creation and destructions

---

<code>pthread_create</code>	creates a new thread
<code>pthread_join</code>	waits for another thread to return
<code>pthread_exit</code>	terminates the calling thread
<code>pthread_self</code>	gets a thread's own ID

### Creating and destroying locks

---

<code>pthread_mutex_lock</code>	locking
<code>pthread_mutex_unlock</code>	unlocking
<code>pthread_mutex_trylock</code>	check lock

Compiling

## *POSIX threads*

Include the header file `#include<pthread.h>`

`gcc fn.c -o fn -lpthread` (other flags as needed)

## *Creating a thread*

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*thread_fun)(void *), void *data);
```

- `pthread_create` creates a new thread of control. It starts execution by invoking function `thread_fun`.
- `pthread_t` is a type (a handle) that gives means for referencing threads. A `pthread_t` variable must exist for every thread created.
- `thread_fun` function executed by `thread`. It must return `void *` and take a single `void *` argument.
- `attr` attributes of `thread` of type `pthread_attr`, default to `NULL`
- `data` argument for `thread_fun` It is passed by reference as a pointer cast of type `void`. If no argument are passed `NULL` is used.
- `thread` terminates by calling `pthread_exit(tNULL)`



## *Creating a thread*

Typical use

```
err = pthread_create(&thread_id, NULL  
                    thread_fun, fun_arg);
```

**err** is zero if success and nonzero if not

## *Creating threads*

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS    8           /* eight threads will be created */

void *PrintHello(void *thr_id)     /* function executed by threads */
{                                  /* can have only one argument */
    long tid = (long) thr_id;      /* user's thread ID */
    printf("Hello World from thread %ld!\n", tid);
    pthread_exit(NULL);           /* OS terminates the thread */
}

int main (int argc, char *argv[]) {
    pthread_t thrs[NUM_THREADS];    /* thread handles are defined */
    int err;
    long t;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %ld\n", t);
        err = pthread_create(&thrs[t], NULL, PrintHello, (void *)t);
    }                               /* no error if err = 0 */
    pthread_exit(NULL);
}
```

## *Creating threads*

```
In main: creating thread 0
In main: creating thread 1
In main: creating thread 2
Hello World! It's me, thread #0!
In main: creating thread 3
Hello World! It's me, thread #2!
Hello World! It's me, thread #1!
In main: creating thread 4
Hello World! It's me, thread #3!
In main: creating thread 5
Hello World! It's me, thread #4!
In main: creating thread 6
In main: creating thread 7
Hello World! It's me, thread #6!
Hello World! It's me, thread #5!
Hello World! It's me, thread #7!
```

## *Passing an entry from global array*

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 3

char *text[NUM_THREADS];      /* defined in global shared memory */
int taskids[NUM_THREADS];

void *PrintHello(void *thr_id) {
    int task_id = *((int*) thr_id);
    printf("Thread %d: %s\n", task_id, text[task_id]);
    pthread_exit(NULL);
}
```

## *Passing an entry from global array*

```
int main(int argc, char *argv[]) {
pthread_t threads[NUM_THREADS];
int rc, t;

text[0] = "English"; text[1] = "French"; text[2] = "Spanish";

for(t=0; t<NUM_THREADS; t++) {
    printf("Creating thread %d\n", t);
    taskids[t] = t;
    rc = pthread_create(&threads[t], NULL, PrintHello,
                        (void *) &taskids[t]);
}
pthread_exit(NULL);
}
```

## *Passing an entry from global array*

```
Creating thread 0  
Creating thread 1  
Thread 0: English  
Creating thread 2  
Thread 1: French  
Thread 2: Spanish
```

## *Passing arguments*

`pthread_create()` permits one argument.

When multiple arguments are needed, a pointer to a structure is passed.

The structure is passed by reference and cast to `(void *)`

## *Passing multiple arguments in structure*

```
#include ....
#define NUM_THREADS 3

char *text[NUM_THREADS];
struct thread_data {
    int  thr_id, sum;
    char *text;
};
struct thread_data thread_data_array[NUM_THREADS];

void *PrintHello(void *thr_arg) {
    int taskid, sum;
    char *hello_msg;
    struct thread_data *my_data;
    my_data = (struct thread_data *) thr_arg;
    taskid   = (*my_data).thr_id;    /* same as my_data->thr_id; */
    sum      = (*my_data).sum;
    hello_msg = (*my_data).text;
    printf("Thread %d: %s  Sum=%d\n", taskid, hello_msg, sum);
    pthread_exit(NULL);
}
```



## *Passing multiple arguments in structure*

```
int main(int argc, char *argv[]) {
pthread_t threads[NUM_THREADS];
int *taskids[NUM_THREADS];
int rc, t, sum;

sum = 0;
text[0] = "English"; text[1] = "French"; text[2] = "Spanish";

for(t=0;t<NUM_THREADS;t++) {
    sum = t*t;                                /* some random calculation */
    thread_data_array[t].thr_id = t;
    thread_data_array[t].sum     = sum;
    thread_data_array[t].text    = text[t];
    printf("Creating thread %d\n", t);
    pthread_create(&threads[t], NULL, PrintHello, (void *)
        &thread_data_array[t]);
}
pthread_exit(NULL);
}
```

## *Passing multiple arguments in structure*

```
Creating thread 0  
Creating thread 1  
Thread 0: English!  Sum=0  
Creating thread 2  
Thread 1: French!   Sum=1  
Creating thread 3  
Thread 2: Spanishe  Sum=3  
Thread 3: German!   Sum=6
```

## *Joining threads*

Joining threads is one way to accomplish synchronization between threads.

`pthread_join(thr_id, status)` blocks the calling thread until thread `thr_id` terminates.

The target thread's `status` is specified in the target thread's call to `pthread_exit()`.

A joining thread can match exactly one `pthread_join()` call.

## *Joining threads*

Steps:

- Declare a variable `attr` of `pthread_attr_t` type
- Initialize `attr` with `pthread_attr_init(&attr)`
- Set the status with `pthread_attr_setdetachstate()`,  
`PTHREAD_CREATE_JOINABLE` (The other option is  
`PTHREAD_CREATE_DETACHED`).
- When done, free the attribute with `pthread_attr_destroy()`

## *Joining threads*

```
#include .....  
#define NUM_THREADS 4  
  
void *SomeWork(void *thr_id) {  
    int i;                                /* loop work */  
    long tid = (long) thr_id;             /* thread ID */  
    double sum = 0.0;  
    printf("Thread %ld starting...\n",tid);  
    srand48(tid);                         /* set seed to tid */  
    for (i=0; i<10000; i++) {  
        sum = sum + drand48()/RAND_MAX;  
    }  
    printf("Thread %ld done. Result = %e\n",tid, sum);  
    pthread_exit((void*) thr_id);  
}
```

## *Joining threads*

```
int main (int argc, char *argv[]) {
    pthread_t thread[NUM_THREADS];
    pthread_attr_t attr;
    int rc; long t;
    void *status;

    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);

    for(t=0; t<NUM_THREADS; t++) {
        rc = pthread_create(&thread[t], &attr, SomeWork, (void *)t);
    }
    pthread_attr_destroy(&attr);      /* free attribute and wait for others */
    for(t=0; t<NUM_THREADS; t++) {
        rc = pthread_join(thread[t], &status);
        printf("Main completed join with thread %ld, status =
                %ld\n",t,(long)status);
    }
    printf("Main thread: program completed. Exiting.\n");
    pthread_exit(NULL);
}
```

## *Locks*

A **mutex** is a lock protecting access to shared data in order to prevent **race conditions**. Procedure:

- Create and initialize a mutex variable
- Several threads attempt to lock the mutex
- Only one of several threads succeeds in locking the mutex
- The owner thread performs some set of actions on locked data
- The owner unlocks the mutex
- Another thread acquires the mutex and repeats the process
- When all are done, the mutex is destroyed

## *Pthreads - mutual exclusion*

(1) Lock must be declared and initialized

```
pthread_mutex_t mutexsum;  
pthread_mutex_init(&mutexsum, NULL);           /* dynamic */  
pthread_mutex_t mymutex = PTHREAD_MUTEX_INITIALIZER; /* static */
```

(2) Acquire a lock on the specified mutex variable `mutex`

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

If `mutex` is already locked the thread blocks until `mutex` is unlocked

(3) Unlock a mutex by the owning thread

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

Error occurs if the mutex was already unlocked

(4) Free a mutex object which is no longer needed

```
pthread_mutex_destroy()
```



## *Pthreads - dot product*

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

#define NUM_THRDS 4
#define VEC_LEN 2000
#define MAX_LEN NUM_THRDS*VEC_LEN

typedef struct {
    double *a, *b;
    double sum;
    int     veclen;
} dot_data;

dot_data dotstr;
pthread_t thr_id[NUM_THRDS];
pthread_mutex_t mutexsum;
double A[MAX_LEN], B[MAX_LEN];
```

## *Pthreads - dot product*

```
void *dotprod(void *arg) {
    int i, start, end, len ;
    long offset = (long) arg;
    double *x, *y, mysum = 0.0;
    len = dotstr.vecilen;
    start = offset*len;
    end    = start + len -1;
    x = dotstr.a;
    y = dotstr.b;
    for (i=start; i<end ; i++) {
        mysum += (x[i] * y[i]);
    }
    pthread_mutex_lock (&mutexsum);
    dotstr.sum += mysum;
    pthread_mutex_unlock (&mutexsum);

    pthread_exit(NULL);
}
```

## *Pthreads - dot product*

```
int main (int argc, char *argv[]) {
    struct timespec start,finish;
    int rc, ntime, stime;
    long i,j;
    void *status;
    pthread_attr_t attr;

    for (i=0; i<MAX_LEN; i++) {
        A[i]=1.1/((double) (i+2));      /* random data */
        B[i]= (double) rand()/100000.0; /* random data */
    }

    dotstr.vecLen = VEC_LEN;
    dotstr.a = A;
    dotstr.b = B;
    dotstr.sum = 0.0;

    pthread_mutex_init(&mutexsum, NULL);
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
```

## *Pthreads - dot product*

```
clock_gettime(CLOCK_REALTIME,&start);
for(i=0;i<NUM_THRDS;i++) {
    rc = pthread_create(&thr_id[i], &attr, dotprod, (void *)i);
}
pthread_attr_destroy(&attr);

for(i=0;i<NUM_THRDS;i++) {
    pthread_join(thr_id[i], &status);
}

clock_gettime(CLOCK_REALTIME,&finish);
ntime = finish.tv_nsec - start.tv_nsec;
stime = (int) finish.tv_sec - (int) start.tv_sec;
printf("main(): Created %ld threads. Time %ld, nsec %ld\n", NUM_THRDS, stime, ntime);

printf ("Sum =  %f \n", dotstr.sum);
pthread_mutex_destroy(&mutexsum);
pthread_exit(NULL);
}
```

## *Other constructs*

Mutexes implement synchronization by controlling thread access to data.

Other synchronizations constructs are

**Barrier**

**Semaphore**

**Condition variables** allow threads to synchronize based on the actual value of data.

## *Pthreads - barriers*

Barrier - a point where the thread must wait for other threads

- the thread will proceed only when predefined number of threads reach the same barrier in their respective programs.

`pthread_barrier_t` - barrier type

`pthread_barrier_init` - barrier initialization

```
int pthread_barrier_init(pthread_barrier_t *barrier,  
    const pthread_barrierattr_t *attr, unsigned int count);
```

- `barrier` - variable used for the barrier
- `attr` - attributes for the barrier, default NULL
- `count` - number of threads which must `pthread_barrier_wait` on this barrier before the threads can proceed further (not specified which ones)..

## *Pthreads - barriers*

Once the barrier is created, each thread will call

```
pthread_barrier_wait()
```

to indicate that it has completed.

```
int pthread_barrier_wait (pthread_barrier_t *barrier);
```

When a thread calls `pthread_barrier_wait()`, it blocks until the number of threads specified initially in

```
pthread_barrier_init()
```

function have called

```
pthread_barrier_wait()
```

All these threads unblock at the same time.

## *Pthreads - barriers*

The main can also call `pthread_barrier_wait()`.

```
int pthread_barrier_destroy(pthread_barrier_t *barrier);
```

The barrier should be destroyed only when no thread is executing a wait on the barrier.



## *Barrier - example*

```
#include ....

pthread_barrier_t my_barrier;
.....
void *thread1() {
    sleep(2);
    printf("Enter integer value for t1: ");
    scanf("%d",&t1);
    pthread_barrier_wait(&my_barrier);
    printf("\nvalues entered:  %d %d %d %d \n",t1,t2,t3,t4);
}
void *thread2() {
    sleep(3);
    printf("Enter integer value for t1: ");
    scanf("%d",&t1);
    pthread_barrier_wait(&my_barrier);
    printf("\nvalues entered:  %d %d %d %d \n",t1,t2,t3,t4);
}
```

## *Barrier example*

```
main() {  
    pthread_t thread_id_1,thread_id_2;  
    pthread_attr_t attr;  
    .....  
    pthread_attr_init(&attr);  
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);  
    pthread_barrier_init(&my_barrier,NULL,2);  
  
    pthread_create(&thread_id_1,NULL,&thread1,NULL);  
    pthread_create(&thread_id_2,NULL,&thread2,NULL);  
  
    pthread_join(thread_id_1,NULL);  
    pthread_join(thread_id_2,NULL);  
  
    pthread_barrier_destroy(&my_barrier);  
    pthread_exit(NULL);  
}
```

## *Pthreads - barriers*

Barriers are for global synchronization

- common when running multiple copies of the same function (SIMD)
- a thread reaching a barrier stalls until all other participating threads reach the barrier.

Example: Iterative linear system solvers:  $\mathbf{x}_{\text{new}} = \mathbf{A} \cdot \mathbf{x}_{\text{old}} + \mathbf{r}$ ;

```
x_new[i*k:i*(k+1)-1] = A[i*k:i*(k+1)-1,:]*x_old;
barrier_wait;
x_old[i*k:i*(k+1)-1] = x_new [i*k:i*(k+1)-1];
barrier_wait;
```

## *Pthreads - barrier*

Not all implementations provide **barrier** functionality.

A crude way to implement barrier with locks


```
barrier(int *counter, int thread_count)
{
    pthread_mutex_lock(&counter_lock);
    *counter++;
    pthread_mutex_unlock(&counter_lock);
    while (*counter < thread_count) {};
}
```

Use a counter that keeps track of how many threads have reached the barrier.

On reaching the barrier, a thread increments the counter and (busy) waits for the counter to reach the value **thread\_count**.

## *Pthreads - hot spots*

Barrier is a **hot spot**



## *Barrier tree*

Spread hot spots, use a tree

